

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model

Creating custom dataset

In [2]: # Please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                           n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_classification.html) for more details

In [3]: X.shape, y.shape

Out[3]: ((50000, 15), (50000,))

Splitting data into train and test

In [4]: #Please don't change random state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)

In [5]: # Standardizing the data.
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

In [6]: X_train.shape, y_train.shape, X_test.shape, y_test.shape

Out[6]: ((37500, 15), (37500,)), (12500, 15), (12500,))
```

SGD classifier

```
In [7]: # alpha : float
# Constant that multiplies the regularization term.
# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.
clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15,
                                penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
clf
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html)

Out[7]: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
                      eta0=0.0001, fit_intercept=True, l1_ratio=0.15,
                      learning_rate='constant', loss='log', max_iter=None, n_iter=None,
                      n_jobs=1, penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                      tol=0.001, verbose=2, warm_start=False)

In [8]: clf.fit(X=X_train, y=y_train) # fitting our model

-- Epoch 1
Norm: 0.70, NNZs: 15, Bias: -0.499391, T: 37500, Avg. loss: 0.552631
Total training time: 0.02 seconds.
-- Epoch 2
Norm: 1.04, NNZs: 15, Bias: -0.750277, T: 75000, Avg. loss: 0.448128
Total training time: 0.04 seconds.
-- Epoch 3
Norm: 1.26, NNZs: 15, Bias: -0.902777, T: 112500, Avg. loss: 0.415726
Total training time: 0.04 seconds.
-- Epoch 4
Norm: 1.42, NNZs: 15, Bias: -1.003874, T: 150000, Avg. loss: 0.400898
Total training time: 0.05 seconds.
-- Epoch 5
Norm: 1.55, NNZs: 15, Bias: -1.075094, T: 187500, Avg. loss: 0.392871
Total training time: 0.06 seconds.
-- Epoch 6
Norm: 1.65, NNZs: 15, Bias: -1.120728, T: 225000, Avg. loss: 0.388065
Total training time: 0.06 seconds.
-- Epoch 7
Norm: 1.73, NNZs: 15, Bias: -1.169943, T: 262500, Avg. loss: 0.385063
Total training time: 0.07 seconds.
-- Epoch 8
Norm: 1.80, NNZs: 15, Bias: -1.203552, T: 300000, Avg. loss: 0.383058
Total training time: 0.08 seconds.
-- Epoch 9
Norm: 1.86, NNZs: 15, Bias: -1.230411, T: 337500, Avg. loss: 0.381694
Total training time: 0.08 seconds.
-- Epoch 10
Norm: 1.91, NNZs: 15, Bias: -1.249466, T: 375000, Avg. loss: 0.380756
Total training time: 0.09 seconds.
Convergence after 10 epochs took 0.09 seconds

Out[8]: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
                      eta0=0.0001, fit_intercept=True, l1_ratio=0.15,
                      learning_rate='constant', loss='log', max_iter=None, n_iter=None,
                      n_jobs=1, penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                      tol=0.001, verbose=2, warm_start=False)

In [9]: clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term

Out[9]: (array([[ -0.83187476,  -0.58935497,  -0.05233948,  -0.59159475,  -0.33771644,
    0.87826212,  -0.85798961,  -0.06890856,  -0.37968271,  -0.3720168 ,
    0.22881296,  -0.04398642,  -0.08060734,  -0.51274272,  -0.07080401]]),
 (1, 15),
 array([-1.24946622]))

# This is formatted as code
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

- We will be giving you some functions, please write code in that functions only.
- After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in `def initialize_weights()`)
- Create a loss function (Write your code in `def logloss()`)
$$\text{logloss} = -1 * \frac{1}{n} \sum_{i=1}^n \text{foreach } Y_i, Y_{pred} (Y_i \log_{10}(Y_{pred}) + (1 - Y_i) \log_{10}(1 - Y_{pred}))$$
- for each epoch:
 - for each batch of data points in train (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in `def gradient_dw()`)
$$dw^{(i)} = x_n(y_n - \sigma((w^{(i)})^T x_n + b')) - \frac{1}{N} w^{(i)}$$
 - Calculate the gradient of the intercept (write your code in `def gradient_db()`) [check this](#)
$$db^{(i)} = y_n - \sigma((w^{(i)})^T x_n + b')$$
 - Update weights and intercept (check the equation number 32 in the above mentioned pdf):
$$w^{(i+1)} \leftarrow w^{(i)} + \alpha (dw^{(i)})$$

$$b^{(i+1)} \leftarrow b^{(i)} + \alpha (db^{(i)})$$
 - calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
 - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
 - append this loss in the list (it will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

```
In [10]: def initialize_weights(dim):
'''In this function, we will initialize our weights and bias'''
#initialize the weights to zeros array of (1,dim) dimensions
#you use zeros_like function to initialize zero, check this link https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
#initialize bias to zero
w = np.zeros_like(dim)
b = 0
return w,b
```

```
In [11]: dim=X_train[0]
w,b = initialize_weights(dim)
print('w =',w)
print('b =',str(b))

w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function - 1

```
In [12]: dim=X_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
return True
grader_weights(w,b)

Out[12]: True
```

Compute sigmoid

$\text{sigmoid}(z) = 1/(1 + \exp(-z))$

```
In [13]: import math
def sigmoid(z):
'''In this function, we will return sigmoid of z'''
# compute sigmoid(z) and return
sigmoid=1/(1+math.exp(-z))
return sigmoid
```

Grader function - 2

```
In [14]: def grader_sigmoid(z):
val=sigmoid(z)
assert(val==0.8807970779778823)
return True
grader_sigmoid(z)

Out[14]: True
```

Compute loss

$\text{logloss} = -1 * \frac{1}{n} \sum_{i=1}^n \text{foreach } Y_i, Y_{pred} (Y_i \log_{10}(Y_{pred}) + (1 - Y_i) \log_{10}(1 - Y_{pred}))$

```
In [15]: def logloss(y_true,y_pred):
'''In this function, we will compute log loss '''
n=len(y_true)
s=0
for i in range(n):
s += (y_true[i] * np.log10(y_pred[i]))+(1-y_true[i] * np.log10(1-y_pred[i]))
loss= (-1 * s)/n
return loss
```

Grader function - 3

```
In [16]: def grader_logloss(true,pred):
loss=logloss(true,pred)
assert(loss==0.07644900402910389)
return True
true=[1,1,1,0]
pred=[0.9,0.8,0.8,0.1,0.8,0.2]
grader_logloss(true,pred)

Out[16]: True
```

Compute gradient w.r.to 'w'

$dw^{(i)} = x_n(y_n - \sigma((w^{(i)})^T x_n + b')) - \frac{1}{N} w^{(i)}$

```
In [17]: def gradient_dw(x,y,w,b,alpha,N):
'''In this function, we will compute the gradient w.r.to w '''
dw = x*(y-sigmoid(np.dot(w.T,x)+b)) - ((alpha/N)*w)
return dw
```

Grader function - 4

```
In [18]: def grader_dw(x,y,w,b,alpha,N):
grad_dw=gradient_dw(x,y,w,b,alpha,N)
assert(np.sum(grad_dw)==2.613689585)
return True
grad_x=np.array([-2.07864835, 3.31604252, -0.79104357, -3.87045546, -1.14783286,
-2.81434437, -0.86771071, -0.04073207, 0.84827078, 1.99451725,
3.67152472, 0.81451875, 2.01062888, 0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)

Out[18]: True
```

Compute gradient w.r.to 'b'

$db^{(i)} = y_n - \sigma((w^{(i)})^T x_n + b')$

```
In [19]: def gradient_db(x,y,w,b):
'''In this function, we will compute gradient w.r.to b '''
db = y - sigmoid(np.dot(w.T,x) + b)
return db
```

Grader function - 5

```
In [20]: def grader_db(x,y,w,b):
grad_db=gradient_db(x,y,w,b)
assert(grad_db== -0.5)
return True
grad_x=np.array([-2.07864835, 3.31604252, -0.79104357, -3.87045546, -1.14783286,
-2.81434437, -0.86771071, -0.04073207, 0.84827078, 1.99451725,
3.67152472, 0.81451875, 2.01062888, 0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_db(grad_x,grad_y,grad_w,grad_b)

Out[20]: True
```

Implementing logistic regression

```
In [21]: def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
'''In this function, we will implement logistic regression'''
#eta0 is learning rate
#implement the code as follows
# initialize the weights (call the initialize_weights(X_train[0]) function)
# for every epoch
# for every data point(x_train,y_train)
#compute gradient w.r.to w (call the gradient_dw() function)
#compute gradient w.r.to b (call the gradient_db() function)
#update w, b
# predict the output of x_train(for all data points in X_train) using w,b
#compute the loss between predicted and actual values (call the loss function)
# store all the train loss values in a list
# predict the output of x_test(for all data points in X_test) using w,b
#compute the loss between predicted and actual values (call the loss function)
# store all the test loss values in a list
# you can also compare previous loss and current loss, if loss is not updating then
stop the process and return w,b
w,b = initialize_weights(dim)
train_loss=[]
test_loss=[]

for i in range(epochs):
pred_train=[]
pred_test=[]
for j in range(N):
dw= gradient_dw(X_train[j],y_train[j],w,b,alpha,N)
db= gradient_db(X_train[j],y_train[j],w,b)
w = w + (eta0 * dw)
b = b + (eta0 * db)
for k in range(N):
pred_train.append(sigmoid(np.dot(w,X_train[k])+b))
LL=logloss(y_train,pred_train)
train_loss.append(LL)
for m in range(len(X_test)):
pred_test.append(sigmoid(np.dot(w,X_test[m])+b))
LL=logloss(y_test,pred_test)
test_loss.append(LL)
return w,b,train_loss,test_loss

In [22]: alpha=0.0001
eta0=0.0001
N=len(X_train)
epochs=10
w,b,train_log_loss,test_log_loss=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```

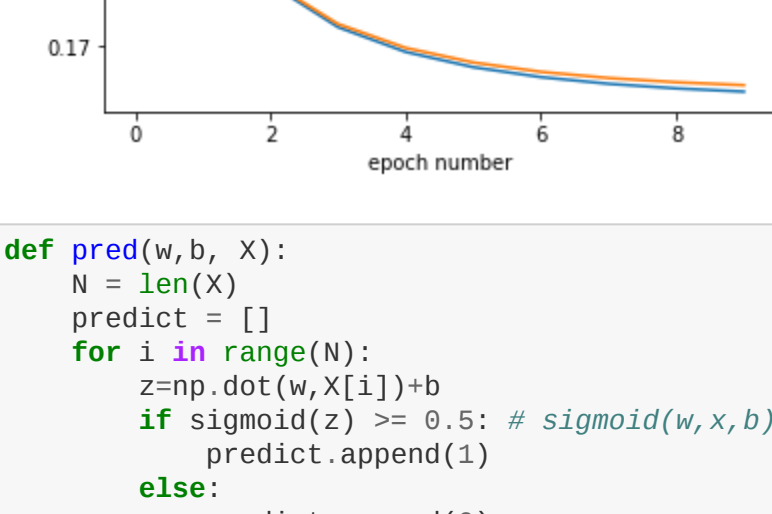
Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

Plot epoch number vs train , test loss

- epoch number on X-axis
- loss on Y-axis

```
In [24]: from matplotlib import pyplot as plt
plt.plot(train_log_loss,label='train_log_loss')
plt.plot(test_log_loss,label='test_log_loss')
plt.xlabel('epoch number')
plt.ylabel('log_loss')
plt.legend()
plt.show()
```



```
In [25]: def pred(w,b,X):
N = len(X)
predict = []
for i in range(N):
z=np.dot(w,X[i])+b
if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
predict.append(1)
else:
predict.append(0)
return np.array(predict)
print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))

0.94944
0.94648
```

```
In [26]: for v in range(epochs):
train_loss=train_log_loss[v]
test_loss=test_log_loss[v]
print("Epochs",v)
print("\ntrain_loss",train_log_loss[v],'\n\t','test_loss',test_log_loss[v],'\n')
```

```
Epochs 0
train_loss 0.20729781784140838 test_loss 0.2072221978118188

Epochs 1
train_loss 0.18556210141426166 test_loss 0.18565259434678275

Epochs 2
train_loss 0.17659652085620509 test_loss 0.17682567720849304

Epochs 3
train_loss 0.17201289496451905 test_loss 0.17235324848189568

Epochs 4
train_loss 0.16938000866115878 test_loss 0.16981009840800462

Epochs 5
train_loss 0.16775336575455 test_loss 0.16825663498220056

Epochs 6
train_loss 0.16669776297615663 test_loss 0.1672612889069227

Epochs 7
train_loss 0.16598837500432867 test_loss 0.16660192986644845

Epochs 8
train_loss 0.1654991822760498 test_loss 0.1661545712175774

Epochs 9
train_loss 0.16515513945496194 test_loss 0.16584572669386236
```

```
In [27]: print('weights:',w,'\n')
print('intercept:',b)
weights: [-0.83543917 0.59554963 -0.05042441 0.59139957 -0.33675676 0.87822749
-0.85715792 -0.068385 0.38429384 -0.36811606 0.22870443 0.6472269
-0.07976159 0.513197 0.07063889]

intercept: -1.247904496218835
```

```
In [33]: #Results we get after comparing our implementation and SGD classifier
print('weight difference:',w-clf.coef_)
print('Intercept difference:',b-clf.intercept_)

weight difference: [[ -3.56441269e-03  6.19465370e-03  1.91566968e-03 -1.95178567e-04
 9.59673686e-04 3.46279712e-05  8.31685290e-04 3.23557144e-04
 4.1113273e-03 -3.90074267e-03 -1.08531791e-04 3.24047785e-03
 8.45749627e-04 4.56977319e-04 -1.65124455e-04]]
Intercept difference: [0.00156173]
```

After implementing logistic regression with SGD and comparing with sklearn's SGD classifier, the weights and the intercept are similar to each other that is within a difference of 10^{-3}